AIFS0023

Written evidence submitted by Phoenix Group

About us

Phoenix Group (Phoenix) is the UK's largest long-term savings and retirement business. We are driven by our purpose of helping people secure a life of possibilities. We passionately believe that our customers are at the heart of what we do and we continue to improve customer outcomes, broadening the focus of our business so that we can help customers as they journey to and through retirement.

Executive Summary

Within our sector, we believe AI will be used mainly to improve customer outcomes and productivity. The ongoing policy development of customer support services like targeted support and pensions dashboards will make obtaining customer insights much more important in the next decade. The government should fully utilise these opportunities and encourage the industry to innovate with increased customer touch points and data, including through machine learning and AI-powered decision support in the future. Developing a clear, future-proof and outcome-driven regulatory framework for both initiatives will be crucial. Currently we see gaps especially in the pensions dashboards regulations, and would ask the government to take an innovation-focused approach and prioritise their development.

In a more general sense, we currently see use cases of AI in analysing investment opportunities, compliance and legal work as well as assisting colleagues on day-to-day administration work. Trials that we have conducted internally have demonstrate good impact on productivity and we have planned further trials with AI assistants.

We believe most risks from using AI could be mitigated with appropriate human intervention and a clear management framework. There are existing rules like Consumer Duty that will require firms to put in place enough safeguarding for consumers, but there might be occasions whether faulty conclusions made by AI originate from poor data or biases that are harder to notice. It will be important for staff deploying AI, especially those involving use of personal data, fully understand the risks and potential harm that their AI tools could bring and intervene in a timely manner. Trainings and human-in-loop processes will be important to ensure positive outcome for consumers.

How is AI currently used in different sectors of financial services and how is this likely to change over the next ten years?

We believe AI will have an impact on both productivity and customer outcomes. Common use cases across these two areas will mostly be related to:

- Assist: AI tools that provide assistant-type experiences, such as Copilot and ChatGPT
- Automate: AI Tools that automate or simplify existing processes
- Analyse: AI that performs enhanced and predictive analytics

Key trends and developments specific to the adoption of AI in the long-term savings industry include:

<u>AI and customer communications</u>

We expect AI could play a big role in customer communications in the next ten years. This could involve generating more tailored comms materials ranging from generating images, videos and explainers. For example, our consumer brand Standard Life currently has a simple chatbot live for

customers and is planning an iterative approach to develop this into a generative-AI assistant for customers. We intend to follow an Alpha-Beta-Live approach with human-in-the-loop safeguards as we progress.

AI could potentially also be used to scale up decision support services and make them more affordable. 85% financial advisers anticipate AI to be used in the advice process soon and majority expect this to happen in the next 2-5 years[1]. Using technology and algorithms to facilitate decision support is not new to the industry, as robo-advisors in the market commonly use algorithms to automate design of investment portfolio for customers and are usually offered at a lower cost than traditional financial advice. AI could play a further role in improving these services and further facilitate the provision of affordable advice. However, we do not expect AI-powered solutions can completely replace human advisers. According to our polling conducted with Ipsos MORI (representative of UK adult population), 44% of respondents said they will not use AI solution instead of paying for an adviser. It should be noted that 45% Baby Boomers (59-80 years old) strongly disregarded the option compared to only 7% of Gen Z (18-26 years old) saying so. While the acceptance for AI-powered advice could change among future generations of retirees, in the meantime it will likely only be playing an assisting role.

While only 8% of customers use financial advice in a year[2], demand for financial decision support is high and FCA's ongoing advice-guidance boundary review will make them much more accessible the next decade. The regulator is currently introducing a new category of regulated support, targeted support, to fill the gap between impartial guidance and full financial advice. The proposal aims to allow firms including pension providers to give more personalised recommendations that provide better outcomes, at scale, for the mass-market of consumers. These recommendations will be provided to a specific segment of customer at specific scenarios, explaining what other people in that situation might do so they could make more informed decisions. FCA expect this support will be free at point of use and based on limited data collected from the customers only. To deliver the support at scale, we expect machine learning will likely be useful in analysing existing and new data, as well as developing new scenarios and segments to support as many customers as possible. We will further explain some of our ongoing work on vulnerability monitoring in later sections.

Pensions data

The success of any digital transformation and innovation using AI, including targeted support, lies in the availability of transparent and good quality data. At the moment, the Long-Term Savings industry is going through a once-in-a-generation customer data audit and digitisation due to Pensions Dashboards. Most pension providers are required to provide data in compliance with the Pensions Schemes Act 2021 and FCA PDCOB in a standard format and connected to the Government-run central architecture by the end of 2025. The government and FCA will then decide when the government dashboard and regulated commercial dashboards could be made available to the public respectively. Once both data and dashboards connection are complete, customers will be able to see all their pensions in a dashboard of their choice. The fundamental change in data availability will likely lead to a proliferation of innovation in the industry, including those powered by AI.

Phoenix Group has been a supporter of the programme since day one. We took part in the Government's Pensions Dashboards Programme (PDP)'s Alpha testing and are the first UK pension provider committing to offer a commercial dashboard. We anticipate commercial dashboards will be able to deliver more than the government dashboard by enabling onward journeys such as targeted support, lifetime modelling or consolidation. AI will likely be used primarily to support these 'post

[1] Schroders Adviser Pulse Survey 2024
[2] Feedback: Advice Guidance Boundary Review | FCA

view services', i.e. services that aggregate data on dashboards so it makes sense to users. While the FCA has already set out their regulatory approach for these services, PDP have yet built the capability for commercial dashboards to connect with its central architecture, not to mention how 'post-view services' will be supported. We believe this capability is not a good-to-have but a crucial to enable savers to have a modern digital journey that they expect. PDP should prioritise making commercial dashboards available without delay. Once these are properly established, we expect the use of AI in the long-term savings products and services can be integrated with wider developments in the financial sector and will increase exponentially.

On top of supporting pensions dashboards services, there is a wider question of what we could do with improved data availability in the industry. Without good data it is difficult to train AI models and obtain useful insights. Based on recent research that FCA commissioned, many customers have unrealistic expectations on how much data their pensions provider hold about them (e.g. knowledge about wider financial circumstances), and thus the level of support providers could offer as a result[3]. It is in fact currently impossible for providers to know people's income sources in retirement without additional input from the customer. As an industry we have limited insights on any individual's holistic financial wellbeing, not to mention on the larger societal level. Unlike in other countries such as the US, there is no database on people's total retirement savings to facilitate research into longitudinal trends of people's saving patterns in the UK. Voluntary efforts have been made through the Pension Policy Institute (PPI)'s Data Project to collate and share good quality anonymised data[4]. 5 pensions master trust providers took part in the private beta phase coordinated by PPI, who concluded that consistent data standards among employers and providers is needed to facilitate data sharing within the industry. This is a persistent gap in pensions data which could be filled if dashboards data could be reused for other purposes and shared securely. Different from Open Banking, the connection and data standards that PDP designed is for enabling data display on Dashboards, instead of for facilitating data exchange within the wider industry. This limits the data's reusability and interoperability with other data initiatives, and potential for analysis and machine learning. The Government should in long term consider how to leverage the Pensions Dashboards data to enable innovations.

We will further explore AI's impact on general productivity in our answer below.

To what extent can AI improve productivity in financial services?

Use cases and benefits

AI is commonly incorporated into most available corporate tool set products in the market. We expect most firms will be adopting these corporate tools where it makes commercial sense and fall within risk appetite. For example, we are currently planning to roll out generative AI assistant to support colleagues on daily basis and exploring using AI agents for HR and IT helpdesks.

As a financial service firm, we have also started trialling AI in work more specific to our business including:

- An AI tool to identify any missing items in draft compliance documents. Before introducing the tool, one in five draft compliance documents contain missing items required by COBS or FCA rule books and manual reviews were required to identify any gaps. With the new tool, we can identify these easily and edit before passing the documents for manual review.
- AI-powered tools have also been used to improve our security posture and analysing or

[3] NMG consulting (2024): Advice Guidance Boundary Review - Targeted Support for Non-Advised Defined Contribution Pensions

[4] PPI (2025) - Pensions Data Project

drafting contracts
- AI Assistants to increase colleague productivity such as note taking, summarisations or information discovery
- We are actively using AI-powered assistants for software developers and data engineers by automating routine coding tasks, generating boilerplate code and explaining complex code patterns to streamline development workflows. Data engineers specifically, these assistants excel at generating queries, optimising data-pipelines and identifying data quality issues.
- More broadly within Asset Management we have several initiatives looking to leverage document summarisation and generation covering Investment Mandate analysis from a compliance perspective, Credit Risk Scorecard Automation and sentiment analysis of Local Authority meeting minutes.

The above trials have demonstrated improvement in efficiency and there are plans to further extend the use of AI to monitor regulatory changes and speed up implementation within the business.

Constraints and challenges

In terms of AI's impact on the workforce, we expect human involvement will still be crucial instead of being completely replaced by AI. However, it is likely that improved efficiency would lead to reduced headcount. For example, if the use of AI in responding to customer queries is proven successful, this might reduce headcount in customer service type roles. On the other hand, it is also important to highlight the need to provide trainings and skills to the existing workforce to remain competitive and adaptive to the changes in ways of working. In the Bank of England and FCA's recent survey[5], 46% of the respondent financial service firms reported only having 'partial understanding' of the AI technologies that they use. While this is mostly due to outsourcing and thus a lack of need to have complete understanding internally, the knowledge gap might reduce firms' ability to identify new ways of using AI in the future as well as potential risks it brings.

As the regulations around AI's deployment is still in development, we expect the regulatory boundary of how AI could be used in financial services will decide whether there might be additional implementation barriers. It will be an intertwined regulatory landscape given the overlap across multiple government departments and regulators' remits. For example with the amount of data processing there might be specific data-related risks that ICO would like firms to mitigate, on top of conduct regulations set out by FCA. Coordination will be necessary to ensure we have a consistent approach and risk appetite towards the use of AI in financial services.

What are the risks to financial stability arising from AI and how can they be mitigated?

Data security and privacy

The utilisation of AI could lead to increase costs and risks relating to cybersecurity if not managed appropriately. There are ways mitigations that firms should consider, for example setting up an AI red team function to identify, assess and help mitigate security vulnerabilities in AI systems will be important. Additional measures including enhancing encryption, regular auditing of any AI systems used, penetration testing could also be used to mitigate the risk. Third party supplier monitoring will be important to ensure data is processed lawfully throughout the system.

Another critical aspect is how will system developers ensure personal data is processed in a manner that does not affect customer rights and freedoms or leading to potential harm. These fairness impacts need to be assessed in relation to:
- societal biases reflected in the datasets used to train the system

[5] Bank of England (2024) - Artificial intelligence in UK financial services

- societal biases that are either explicitly or implicitly reflected in the training data
- the decisions made by teams during the AI development and deployment lifecycle
- systems can possess characteristics that, while not necessarily reflective of societal biases, can still result in unfair behaviour when these systems interact with stakeholders after deployment
- unfair allocation of opportunities, resources, or information
- failure to provide the same quality of service to some people as they do to others
- over or underrepresent groups of people, or even treating groups as if they don't exist

Finally, threat actors use of AI is also being monitored closely with easier access to social engineering, phishing and automated attack tools likely to emerge in the near term.

We expect financial service firms will be addressing the above in order to comply with the FCA's Consumer Duty.

Data maturity – input and output

Hallucinations and bias are all derived from the inputs and the best approach is to get these right in the first place, although it will still be possible to train out bias in a computer system. Consideration must be given to how the accuracy of the personal data could be compromised and what harms this could cause to relevant individuals. This is particularly relevant in instances where toxic data entered into the system. Reviewing model training data to ensure data input and output meet requirements concerning data quality and accuracy will be crucial.

The key specific risks which AI brings comes with its speed of distribution and confidence. Any errors can be spread far quicker in an AI automated system compared to, for example, a single human operator making errors. Also unlike humans, the models will usually present information with the utmost confidence – often leaving users with no uncertainty to the outcome. The use cases of AI is expected to go beyond current trends for example in credit underwriting and other risk management process, and could lead to unexpected impact on the market and individuals. Given the potential impact, financial service firms should ensure they can explain simply to customers how AI is used to process their personal data through the corresponding systems and how the decisions it makes impact them. In the case that AI processing have to be ceased, it will also be important to understand which part of the process the AI influenced outcome and how the absence of AI might change the level of risk involved in the processing.

Overall, mitigations that financial service firms could be put in place include:

1. Design – quality control of data input when forming models – an important point here is that companies understand how the models are built and the inference it places on different data points
2. Alpha-Beta-Live – avoiding deploying AI on an immediately large-scale; trialling in-house and then with test customers.
3. Ongoing evaluations – tests that are run automatically to allow colleagues to check that the models are performing within their set parameters
4. Manual review – sampling and constant tweaking of models post-launch
5. Maintain human-in-the-loop for processes with high degree of risk decisions
6. Deploy technical guard rails that prevent harmful input and output
7. Develop organisation specific ethical guidelines for AI use
8. Ensure that business subject matter experts have oversight of application specific models and their outputs

What are the benefits and risks to consumers arising from AI, particularly for vulnerable

consumers?

As mentioned in the first section, we expect AI could support us in delivering more tailored communications and support for customers. With regards to vulnerable customers, we are currently exploring using voice analytics to understand customer needs, including positive assurance around identification and response to characteristics of vulnerability. Some other use case could also be surfacing next best actions for customers driven by the data we hold on them and what they tell us during an interaction, enabling much more personalised experiences and communications. We expect AI driven knowledge management will help contact center staff respond accurately and quickly to customer queries. In the future, there could be further developments of innovations in how vulnerable customers could communicate with providers and how support is being provided, e.g. more diverse channels like chatbots and voice assistants. Integration of pensions dashboards and targeted support will be crucial to develop more suitable support, including through automation, to cater different needs.

Application of AI in customer-facing work requires stringent monitoring of the quality of data input and data protection. Although not unique to its application in financial service, communications generated through large language models could amplify unconscious bias. Any services supported by AI will require consistent review to ensure positive consumer outcome and compliance with Consumer Duty. The AI system could affect the rights and freedoms of individuals and cause them adverse harm. For example, AI could be used to individually assess risk, predict health outcomes and refuse or propose unacceptable premium increases for the most vulnerable people due to their characteristics. It is paramount that any relevant demographic groups and other contextual details are accounted to ensure fairness concerns are identified and mitigated at the earliest juncture. Furthermore, the 'Human in-the-Loop Oversight and Intervention' must be part of the processing activity to ensure that the results of a customer application are examined by a Human before the results are returned to the individual.

The extent to which AI influences outcomes without awareness of the user utilising the tool should also need to be considered. There is a key question around the safeguards applied here to the processing of special category data in particular, given that the processing of vulnerable customer data will very likely include such special category data. There could also be instances where customer objects the processing of their personal data via the AI solution due to concerns over its outputs. Alternatively, the customer might request a restriction of processing of their personal data for a specific part of the processing activity involving an element of AI.

If the customer's consent was obtained for AI processing, particularly that of a vulnerable customer, how was it obtained, where is it stored, what were they told at the time and how to withdraw it will be important. If they withdraw consent of for using their personal data via the AI solution due to concerns over its outputs, consideration should be given to how firms can service them to ensure there is no compromise in the level of service provided. This is a critical aspect that can be defined and accounted for during Data Protection by design & default both within the organisation internally and by the third-party AI provider.

How can Government and financial regulators strike the right balance between seizing the opportunities of AI but at the same time protecting consumers and mitigating against any threats to financial stability?

We appreciate the Government's response to the AI regulation white paper, and the FCA's follow up update on their approach to AI. Both documents were encouraging and committed to ensure UK legislation and regulations will remain technology agnostic and principle-based, which we think was reflected in the recent consultation on targeted support and the pensions dashboards regulations.

While it is important that the Government and regulators take an outcome-based approach, in cases where clear public good can be demonstrated, proactive work should be carried out to drive innovation. Guidance, research and public forums should be used to facilitate good practice sharing and maximise delivery of public good. This will be particularly useful for the development of new Government-led programmes like targeted support and pensions dashboards to maximise public-private collaboration.

When designing regulatory frameworks for new innovations, we will also need to be careful of dependents across new policy areas to avoid creating regulatory paradoxes. In the case of Pensions Dashboards, the drafting of design standards for commercial dashboards relies on live data testing, but live data cannot be connected to commercial dashboards until they can comply with regulatory requirements including design standards. While the PDP and MaPS are now testing with the government dashboard as a workaround, this has limited private sector's ability to test their own propositions and will delay the launch of commercial dashboards and post-view services.

Another area that the Government and regulators might want to proactively act on is the knowledge gap in AI technology. As mentioned in the previous section, nearly half of the firms responding using AI said they only have partial understanding of the AI systems they deployed. There needs to be requirements ensuring firms conduct regular audits of the systems they deploy and those who are responsible for auditing are qualified. Qualifying criteria could include taking an industrywide-recognised AI CBT training programme, focusing particularly on the potential risks and harms of AI. Regulators should work with organisations in designing these requirements and how often should they be reviewed.

As FCA's Principles for businesses and Consumer Duty currently sets out, firms must pay due regard to the information needs of its clients and communicate information to them in a way which is clear, fair, and not misleading. Firms should continue to follow existing requirements when explaining to customers how AI is utilised to process their personal data in plain language. Communications should be timely and cover sufficient details when it is in customer's interest.

March 2025