AIFS0071

Written evidence submitted by Deloitte

Deloitte welcomes the opportunity to contribute to this inquiry led by the Treasury Select Committee.

Our Financial Services Practice possesses extensive experience of working with clients of all sizes across the UK, utilising expert market insight and technical expertise to enable development and growth for firms and the wider economy. In short, we are determined to build a better financial system – one which is safer, more sustainable, digital, and accessible.

In partnership with the Deloitte AI Institute, we leverage deep industry knowledge to lead the AI conversation in the financial services sector and uncover insights to make sense of what is a complex ecosystem. Within our own global organisation, we have developed PairD, Deloitte's virtual AI helper, first launched at Deloitte UK in October 2023, which our people are using to accelerate and enhance the work we do for clients (with appropriate controls over client proprietary data as required), internally and with charitable partners, such as Scope.

For the financial services industry, we are clear that to ensure its own growth and contribution to the broader UK economy, it must continuously evolve across its entire ecosystem. As new developments emerge, driven by technological advancements and evolving consumer needs, the sector needs to adapt and innovate to capture emerging opportunities while ensuring protection from evolving risks.

The foundation is never static, and the arrival of Generative AI (GenAI) marks a new era, exploding the number of potential use cases and putting benefits in the hands of the workforce. In a five-year timeline, our financial services team sees GenAI as the single biggest controllable opportunity for FS firms to improve their competitiveness.

But to make this a reality, there will be challenges to address and important questions to answer. We have provided high-level views in this submission, which we would be happy to expand upon in further dialogue with the Committee. As set out below, we have particular interests in:

- The adoption and use trends of AI in the FS sector.
- The extent to which AI can improve productivity.
- Benefits, risks, and guardrails for organisations.
- Effective and proportionate regulation.

Such is the breadth of our firm, our insight stretches across multiple sectors, including investment management, insurance, real estate and more. We will touch upon all of these in our submission, but to provide additional specificity, many use cases will focus on our insight into banking and capital markets.

How is AI currently used in different sectors of financial services and how is this likely to change over the next ten years?

AI is not new to the financial services sector (FS). It has been in use for years in specific functions, including algorithmic trading and trade surveillance. Indeed, traditional AI excels at analysis and automation and has been in use for some time now. But the arrival of GenAI marks a new era, dramatically increasing the number of potential use cases, putting benefits in the hands of the workforce, and creating opportunities for growth in the sector.

This capability has proved to have game-changing potential for meeting the challenges faced by banks and capital markets today, and benefits include:

- **Increased efficiency, and reducing operating costs** is perhaps GenAI's most well-known benefit. Whether it's automating repetitive tasks such as data entry or analysis, processing vast amounts of information with greater precision and fewer errors than humans, producing operational content such as meeting minutes or summaries, or conducting faster searches of complex data using natural language, GenAI is saving financial service firms time, and enabling more effective resource allocation.

- **Enhanced customer experience** is another area in which GenAI promises to deliver high impact. As anyone who has ever opened an investment account can attest, new client onboarding involves a lot of filling out and signing of documents, an arduous process for both financial service institutions and their customers. Once a client is on board, there's still the matter of understanding and managing their assets and identifying the best opportunities for their particular portfolio – an increasingly challenging task as asset classes expand and become more complex. Yet today's consumers, investors, and corporate customers expect a fast and smooth onboarding experience, plus the best advice and asset management available, quickly.

- **GenAI is proving instrumental in making digital agents more personalised**. Today's GenAI-powered agents are summarising conversations intelligently, offering similarly conversational responses, acting with human-like empathy, and answering an increasingly complex range of customer requests. The result has been reduced customer wait times, and less need for human intervention as digital agents learn how to answer more, and more complex, questions. GenAI can take personalisation a step further. With the ability to analyse customer preferences and behaviours, a GenAI-powered digital agent can also recommend financial products and services that are tailored to individual customer needs. Ultimately, that digital agent could customise pricing in real-time, delivering competitive offers to target customers, such as preferential lending rates, based on an enhanced measurement of their credit risk.

- **Enabling understanding is another of GenAI's strengths**. This is particularly valuable for financial service organisations, which are not only information intensive, but often have data stored in multiple locations, in the cloud and within local legacy systems. For banks and capital markets, this GenAI capability represents a double benefit. To understand and predict market trajectories and make prudent buying and selling recommendations, analysts must sift through an array of company filings, transcripts, reports, news, interest rate updates, and risk profiles – a time-consuming process, precisely when speed matters. At the same time, querying multiple databases from multiple locations adds further hurdles to retrieving relevant information quickly.

- **Compliance and security** is one more area in which GenAI is making an impact. The technology's ability to see patterns and notice out of the ordinary activity makes it well suited to fraud detection and loss prevention. That same capability can favourably impact compliance activity as well. When this work is completed by humans, the potential for errors often exists. GenAI, on the other hand, can process repetitive content faster, and with fewer inaccuracies, while also checking things like localised marketing content in different languages for regulatory matters within each jurisdiction – making it a useful tool for a human agent to deploy and oversee.

Across sectors, we are seeing tangible examples of organisations realising these benefits, including fraud reduction and increased click-through for personalised lending on customised offers for consumers in retail banks. In corporate and transaction banking, we are seeing banks fully automate the loans underwriting process; in investment banking (IB) and capital markets, Deloitte US reported in 2023 on the potential of a 27% productivity increase across investment banks and 27% – 35% front office employee productivity by 2026; and in the insurance sector, we have identified reports from

others[1] where underwriting teams at a specialised insurer experienced a 113% productivity increase using GenAI-supported workflows for underwriting submissions relating to bespoke policies.

These are just a few examples, reflective of a broader trend, not just in the FS sector, but the wider economy. Findings from our latest State of Generative AI in the Enterprise[2] study showed that overall, two-thirds of business leaders (67%) reported their organisation is increasing its investment in GenAI. Alongside this, our 2024 Digital Consumer Trends research found that workers are increasing their use of AI too, with 39% of GenAI users claiming they used it for a work purpose[3]. This is equivalent to 14% of the UK, and 7 million people between 16-75. This is up 66% from last year (May 2023), at which point 4 million in the UK had used Gen AI for work.

It is notable, however, that a large proportion, possibly the majority, of those using Gen AI for work purposes may have done so without official endorsement by their employer. A substantial number of companies across the economy (nine in ten) still lack a policy on whether use of Gen AI is acceptable, or a governance structure for if it is.[4]

In a rapidly evolving area, the speed of deployment, including within organisations, will naturally vary. We observe that "banks move at the speed of banks" rather than at the speed of technological advancement. However, at a time when costs for financial organisations are increasing, while profits from traditional income sources are down[5], GenAI presents an answer to help address key challenges for the FS sector.

How is the use of GenAI likely to change over the next ten years?

Looking ahead, given the speed of development, the range of possible use cases are extremely broad and evolving. However, there are trends to explore, including where GenAI is maturing. This includes:

- **General Intelligence:** the mimicry of human thinking will likely mature to include reasoning, making new inferences and predictions based on complex inputs. It will start to form perspectives and views that challenge human thinking. Currently GenAI focusses on distilling, summarising, and producing human-like content.

- **Versatile Applications:** AI will handle increasingly diverse types of input and output and the available combinations of input/output modalities will increase, with greater sophistication in the production of specialist output such as financial engineering constructs, systems architecture design, audit opinions and risk control assessments.

- **Artificial EQ:** interactions will become more human-like, developing the ability to change tone and respond appropriately to human emotions. This will be particularly impactful, making AI suitable for an increasing number of human-facing use cases.

- **Agentic AI:** represents another significant leap in artificial intelligence, moving beyond reactive systems to those that operate with a degree of autonomy and proactivity. An Agentic AI capability does not just respond to commands but acts as a capable "agent" working towards goals. These AI agents draw on multiple types of AI capability, including GenAI (to generate content) but also other forms of AI including Machine Learning, Deep Learning, Natural Language Processing (NLP) and Knowledge Representation & Reasoning (KRR).

---

[1] Cytora. Case study - Markel uses Cytora and achieves +100% productivity uplift to fuel growth. September 2023.
[2] Deloitte. State of Generative AI in the Enterprise Q4 Report, 2025.
[3] Deloitte. Digital Consumer Trends UK, 2024
[4] Deloitte. Digital Consumer Trends UK, 2024

Future use of Agentic AI will represent the next frontier of possibilities for scaled automation and performance improvement within financial services.

Such is the speed of maturation we are seeing across these themes; we believe AI is on the threshold of a paradigm shift. For FS sectors, already we are seeing the transition of GenAI from an 'instrument of strategy' (i.e., accelerating delivery of today's business plan) to a 'determinant of strategy', where tomorrow's business is planned around new AI capabilities.

In the years ahead, the competitive landscape in FS will likely be redrawn, with the sector's winners and losers determined by the speed and effectiveness with which their AI initiatives enable evolution of their business operations, products, and services. As highlighted, it is difficult to extrapolate the entire breadth of use cases in the future.

However, just as significant as the evolving use of GenAI in the next ten years, is the internal impact on organisations. This is an area where our firm possesses significant expertise, in our role supporting business transformation and adoption of AI.

The organisations expected to capture the biggest benefits from prior waves of technology-enabled innovation (e.g., cloud, data, digital) will continue to outperform in driving value from AI. This is because leaders in innovation have already invested in the key organisational enhancements, including culture, governance, data management and agile delivery methods, needed to capitalise on the AI opportunity.

Key foundations will need to be in place for firms, including the technical capability to unlock the benefits of AI. Some of these advancements which have enabled prior progress include:

- **Modernised Data Platforms**, where done well, delivered readily scalable computing power and accessible data provisioning, that abstracted data away from the complexity of legacy architectures while reducing total cost ownership of the IT estate. Where public Cloud solutions were used, it also forced banks to learn how to assess and manage the risks associated with introducing third-party dependencies to the infrastructure supporting core business processes.

- **Data Governance** may have been implemented initially for compliance purposes but has established the organisational accountabilities, policies, quality improvement methods and understanding of organisational data assets to provide trusted datasets as inputs to AI use cases.

- **Digital Banking** has evolved customer expectations to be more comfortable with self-service. Furthermore, key office functions such as client onboarding and loan fulfilment can be more easily fulfilled with AI assistance.

Alongside this, the businesses which have realised benefits from prior technical innovations will have learned and refined the delivery methods that work in their organisation. Typically, these have included consideration of:

- **Governance** – putting in place sufficient oversight to adequately assess and mitigate the spectrum of risks, without unduly constraining delivery.

- **Risk & Control Framework** – embedded into run-the-bank and change-the-bank processes and regularly updated for evolving risks.

- **Culture** – benefits are well communicated; business function owners expect to embrace emerging technology to improve process performance.

- **Idea to value** – strong change management and adoption processes are embedded to generate ideas from innovation, assess feasibility and investment case, to rapidly deliver the best ideas into production and scale.

- **Talent** – hiring and learning/development approaches that build adequate skills and capacity; and,

- **Partnerships** – engaging with the wider market ecosystem, forming partnerships with technology and service providers best placed to assist delivery.

The trends that dominate the years ahead will be informed and shaped by the work to unlock the true value of this technology, and how processes and how banks are reconceived internally. As banks evolve in maturity with GenAI they will also begin to give front line employees increasing autonomy and improved tooling that will enable increasing revenue while also reducing time spent on tasks such as data entry. But once that tooling is in place and banks begin to reconceive processes there must be a focus to continue to redeploy staff to higher value roles. It is important that the continuous upskilling of teams who use these new tools to do more is not a one-time effort, but instead is something built into the talent model and measured.

Certainly, in a five-year timeline, our financial services team sees AI as the single biggest controllable opportunity for organisations to improve their competitiveness.

Adoption of AI in the FinTech Sector

The Committee has asked about AI adoption for FinTech firms. It is reasonable to suggest that the FinTech subsector is likely to move quickest in the adoption of AI, due to distinct execution advantages.

These include the relative simplicity of their current operating models (considering products, processes, technology, data, and organisation) making them less encumbered by the constraints of legacy systems and processes. They also still have the flexibility to jump straight to newly conceived processes without lengthy re-engineering of legacy system architectures and processes.

Furthermore, they also typically have a culture tilted to more rapid growth and innovation. If paired with robust compliance and control systems, this positive risk appetite can mean they will be more willing to extend AI capability to customers and into production processes sooner.

To what extent can AI improve productivity in financial services?

Ultimately, the significance of AI to the sector will be measured by the extent that this innovation delivers sustained impact. Areas where this can be realised include:

- **Increased efficiency** – automate repetitive tasks, freeing human resources for more complex, creative or customer facing engagement.

- **Improved accuracy** – process vast amounts of data with greater precision and fewer errors than humans, leading to more accurate predictions and outcomes.

- **Enhanced personalisation** – analyse customer preferences and behaviours to create tailored experiences, improving customer engagement.

- **Trend prediction** – make data driven decisions, detecting trends, and predicting changes in the market.

- **Creativity** – new possibilities to create new possibilities for products, services and business models fostering innovation and growth.

- **Cost savings** – streamlining operations, reducing errors, and enabling better decision-making, AI can help save costs and allocate resources more effectively.

- **Protection** – improving the effectiveness of financial crime and loss prevention capabilities.

- **Accessibility** – Make the services more accessible and affordable.

At the moment, most banks are currently building AI business cases around cost reduction, and this is no surprise.[6] It is easier to get funding approved for initiatives which drive out cost. The impact tends to be delivered quicker and benefits tend to be more directly attributable to the investment made.

It is worth stressing our view that most benefits come through combining human and AI strengths, not through large scale replacement of humans with AI. We believe institutions should develop and strengthen human skills to allow for adoption and value realisation. In support of this, we have identified three key modes for achieving value through AI, all of which combine AI and human strengths.

We have identified the persona of core "agent modes". We use the word "agent" to indicate a new capability that combines AI and human inputs in its operation.

- **Assist:** A focus on productivity through personal agents will help employees work faster and smarter.

- **Augment:** A focus on improving quality and process performance through Specialist Agents will help experts in areas like risk management or data analysis do their jobs better.

- **Automate:** The possibilities of overhauling processes are increased, using AI to completely redesign how things are done.

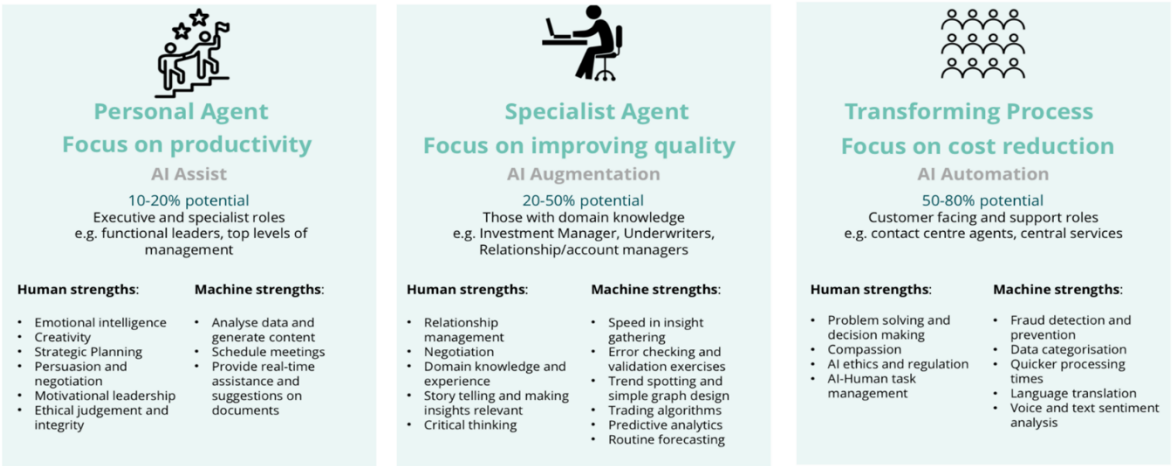

As previously highlighted, cost reduction is the most common business case presented at this time. Themes of this will likely include:

[6] Avanade. "Rethink: How banks can drive down cost and protect profit". 2020.

Workforce acceleration

A “marginal gains” approach to deploying many productivity improvements across the human workforce, making small improvements that add up. At the most basic level, this will include automation of repetitive tasks such as data entry and analysis, search and query, draft production of many varieties of operational content (meeting minutes, communications) and summarising large documentation. This is the type of “text and images” productivity support perhaps most associated with GenAI, particularly among newer users.

However, we see that the art of the possible is rapidly expanding, with more specialist acceleration use cases including data governance and management, data quality and remediation, model development and analytics.

Workforce acceleration will also likely require widespread uplift in workforce skills with AI in the same way as staff previously became proficient in typing, spreadsheets and calendar management and other functions which historically were performed by specialist resources only.

Engineering transformation

Banks typically have large, complex systems landscapes which are highly customised to the specific organisation and may include core components ranging up to 50 years in age. GenAI can generate and optimise software code, reducing the time to write, while improving quality. This can accelerate production releases and make maintenance less onerous, particularly for engineers with less experience of specific bank systems.

Loss avoidance

Risk management, fraud prevention, cyber, legal, and other brand protection functions have high potential for improvement through AI. These functions tend to be improved by speeding up processes, expanding the scope of processes, and providing wider sets of data inputs to improve process performance – all of which AI readily supports humans in. Specifically, AI-enhanced credit risk management improvements can result in fewer loan impairments and write-off charges. Fraud prevention and financial crime (FinCrime) processes can be accelerated and expanded using AI to review a wider set of input data sets to uncover new insights on actors and ultimately reduce losses.

While more difficult than reducing costs, FS business will likely also invest to increase productivity to grow revenue. We see several key revenue-impacting themes, including:

New capabilities for growth

- Insight-driven pricing: real-time customisation of pricing (e.g., preferential lending rates) to make highly competitive offers to target customers based on enhanced measurement of their credit risk.

- Hyper-personalised marketing: improved conversion rates based on insightful identification of individual prospect and customer/client needs, and highly tailored communication.

- Next generation trading algorithms: trading income uplift from enhanced market insight and automated trading decisions

Customer experience and retention

AI-powered digital agents (e.g., chatbots) can reduce customer wait times by addressing an increasing range and complexity of customer requests. While certain customer journeys (e.g., those associated with large transactions, bereavement etc.) must remain as person-to-person interactions, the improved

responsiveness of digital customer service agents can improve customer experience and retention rates. Increasingly the quality of AI interaction with humans has the potential to improve as AI technology develops–adjusting the AI agent's behavior according to the behaviour/emotions of the customer. Of course, it should be noted that whilst there is significant potential here, some customers have reported frustration with being unable to easily access person-to-person services. The onus is rightly on organisations to improve processes and outcomes produced by AI agents and provide meaningful alternative options for the cases it is not appropriate or desirable from the customers perspective.

Stepping back, depending on organisational priorities, i.e., achieving cost efficiencies, risk reduction or focussing on the customer experience and growth agenda, different functions will have differences in potential. Organisations prioritising customer/client experience and a growth agenda will likely naturally focus more on the front office sales, marketing, and service functions, and in tailoring product/service features and pricing. Interestingly, for FS, we see credit risk management as a key enabler of growth, where greater precision in credit-risk assessment enables insight-driven pricing of lending to secure the most desired customer segments.

More broadly, we expect that whichever components of driving value are prioritised, results will come from building and implementing enhanced and/or new capabilities that combine human and AI strengths. Deloitte has defined an initial priority set of over 50 banking-specific use cases, many of which we are already working with clients to implement[7], which will likely have widespread impact across end-to-end banking operations. A selection of examples, including the business challenges they address, are described below. Here we use the word "Agent" to indicate a new capability that combines AI and human inputs in its operation.

Business area	Typical business challenges	AI opportunities
Strategy and proposition	• Increased competition • Specialist resource constraints • Internal siloes across propositions. • Further data required to inform strategy and proposition development	• **Research agents** to summarise and perform trend analysis of market, customer, channels, and products to support strategy and proposition development.
Sales/trading and marketing	• Evolving customer/client landscape and expectations • Identifying individual needs and tailoring products/services accordingly • Omni-channel marketing and personalised content • Financial market volatility	• **Hyper-personalisation communication and personalisation agents** driving personalised messaging and marketing campaigns based on market available data, interaction data (voice, digital), etc. • **Creative content creation (multi-modal)** to accelerate delivery of the marketing content and align to brand, legal and compliance policies.
Customer/client experience	• Areas to improve in customer/client-centricity. • Capturing and managing customer/client feedback	• **Virtual agent scaling** the operations centre automating end-to-end client experience. • **Assistant for channel**

[7] Deloitte. "Changing the game: the impact of artificial intelligence on the banking and capital markets sector", p19. 2024.

	throughout touchpoints • Establishing seamless omni-channel presence integrating physical and digital experiences • Availability of insights to customer/client profile for service agents/risk management (RMs)	**engagement** (RM, call-centre, branches, and email) improving efficiency, driving improved customer experience.
Operations and service	• Straight-through-processing and real-time capability • "One pattern for all" model in legacy systems/processes restrict personalisation. • Areas to improve in data management. • Areas to improve onboarding and service journeys due to organisational/ technical silos.	• **Client / Customer on-boarding** by gathering information, assessing trustworthiness, and seamlessly onboarding them onto new products. • **Customer service** managing complex requests seamlessly, helping customers engage with their bank and providing them with useful responses. • **Fees collection and recovery services** for early detection and helping draft appropriate communications.
Risk and compliance	• Complex and evolving regulatory requirements, particularly for global institutions/organisations. • Emerging risks due to rapid technological change – particularly AI and cyber • Areas to improve in data management. • Balancing risk appetite and growth ambitions	• **Regulatory compliance** by detecting, alerting, and generating triggers for the customers. The summarisation for compliance reporting. • **Financial crime support** activities like, narrative generation. Adverse media screening and other potential fraudulent threats. • **Cyber detection services** to look for any anomalies and detect patterns.
Technology	• Leveraging technology to improve operational efficiency, enhance customer experience, maintain compliance, and ensure security – all while remaining agile and competitive in an evolving financial landscape. • Legacy systems with poor real-time capability • Skills and capability challenges in retaining knowledge of legacy/proprietary systems and new technology	• **Delivery of engineering and operations support** to drive quick delivery of products/solutions to market.
Data and analytics	• Banks are data-intensive organisations but there can be areas to improve in data management and processes. • Stringent regulatory requirements	• **Accelerated insight** production where business can use where natural language to query complex datasets without reliance on

	increase the load on business and chief data officers and their teams. • Increasing dependency on data and analytics to deliver business growth objectives	technical data staff. • **Greater coverage of data governance** standards, where GenAI accelerates data profiling, definition, and quality improvement.
Other enabling and control functions	• Complex financial regulation • Legal litigation and disputes	• **Focus on functions like, HR, finance, legal, etc.** to develop specialist agents to integrate into business processes providing summarisation and other functions.

Ultimately, a significant component of the importance of AI to the sector will be assessed on the extent that this innovation delivers sustained profit and loss impact. Here we consider the "size of the prize" given a typical cost-to-income ratio profile today and expected AI benefit themes. Doing this, we have seen reported a potential 5-7% positive contribution in 2-3 years, and 10-15% in 5-7 years to cost-income ratio.[8] This view considers a wide range of banks, and smaller, more nimble organisations including those with currently high cost-income-ratios (CIRs), who would find greater opportunity to achieve the higher end of this 5-15% range of improvement, including FinTech firms.

What are the risks to financial stability arising from AI and how can they be mitigated?

As well as significant opportunities for efficiency and innovation, GenAI also presents unique challenges. Of course, the financial sector, like any other, is impacted by the global and geopolitical environment in which it operates. The advent of new technology always presents risks and opportunities to the existing order, and these have been discussed extensively, particularly when it comes to accessing the critical infrastructure that will power the GenAI.

Despite this, it should be remembered that to a large extent, financial stability is also dependent on the actions taken by individual institutions. Without organisations putting robust risk management in place, we cannot expect consumers to have appropriate confidence in the trustworthiness, fairness, and stability of institutions – presenting a risk to the wider ecosystem.

A proactive, comprehensive approach that prioritises both robust risk management and the preservation of public trust is therefore essential to harnessing benefits, while safeguarding the integrity of the financial system. Therefore, for the purpose of this inquiry, we will explore this more narrowly defined area, to focus on measures individual organisations can take, which can help preserve stability.

First, for banks and others, numerous risks need to be managed effectively. These include:

- **Misuse of AI**: Incidents involving malicious actors are a real threat. For example, where "deepfake" techniques have been used to successfully imitate a customer, gaining access to their account to commit fraud.

[8] Sourced from "Refinitiv", "Factiva", "Statista" selected Bank Annual Reports as available in Q4 2023.; "Deloitte AI Institute". 2024.; "Gen AI Dossier". 2024. Deloitte. Changing the game: the impact of artificial intelligence on the banking and capital markets sector, P5.

- **Environmental impacts**: Greater use of AI will naturally require greater computing power, in turn leading to greater energy consumption in the data centers. This is a competing factor against industry net zero commitments and climate change impact disclosure and reporting requirements.

- **Amplification of biases**: Underlying datasets contain inherent biases that will be amplified once the models are trained on them, potentially exacerbating any discriminations based on gender, race, and other characteristics. This could expose banks to litigation and/or regulatory consequences.

- **Cyber risk**: The risk of irresponsible application of AI is associated with the various use cases that will likely get contemplated by businesses (e.g., using Large Language Models (LLMs) for heightened automated cyber threats). Additionally, where AI starts to form substantial components of an organisation's operations, lack of transparency into precisely how the AI is functioning could potentially leave the organisation unable to design appropriate controls and cyber protections.

- **Sovereignty restrictions**: The sovereignty risk relates to the expectation that AI models trained on certain data sets are subject to sovereignty/residency regulations and will be required to be run only on data centers within that jurisdiction. Additionally, organisations with global processes, for example trading operations, will face complex challenges in satisfying divergent local regulations across the jurisdictions in which they operate.

- **Future regulation**: The risk of lack of certifications concerns the possibility that LLMs may face future regulation as they are increasingly used for insights, advice, etc.

- **Autonomous vs. human intervention**: The risk of safe usage is associated with how and where LLMs are used (e.g., using LLMs to generate autonomous action for machinery in a factory floor).

- **Fight for talent**: Across the industry, the need for skilled science, technology, engineering, and mathematics (STEM) capability, including data science, is on the rise. Organisations will likely struggle to retain talent and will need to put in place upskilling and recruitment initiatives.

Whilst the Committee will reflect on the full range of options for policymakers to intervene, we do have a view on specific frameworks and guardrails which businesses should consider – both to mitigate risk, but also to ensure the benefits of AI can be realised.

This includes through Deloitte's Trustworthy AI framework[9], which sets out guardrails to address key risks. These guardrails include:

- **Secure environments:** Train LLMs in secure environments in a data centre or on cloud to reduce probability of leakage of company information.

- **Restricted usage:** Restrict initial usages of generative AI to increase accuracy of inferences; then scale once there is a growing sense of comfort.

- **Enterprise data sets:** Train LLMs with data sets that are governed within an enterprise rather than the wider internet.

[9] Deloitte: Deloitte's Trustworthy AI framework

- **Audit trail:** Persist the data that LLMs are trained on to trace the data, map the lineage, and have an audit trail of what type of data was used.

- **Trust but verify:** Keep humans involved throughout the process to assist in validating and verifying the generated output and to monitor the AI's accuracy.

- **Large language model operations (LLMOps):** Form dedicated team(s) focused on operating, managing, and governing the models to prevent drift over time and for rooting out biases.

It is important that organisations execute their rollout of AI with a clear vision, and one that addresses concerns for the specific shape of their enterprise. For financial institutions, responsible AI adoption must prioritise robust risk management, data privacy and security to mitigate the associated perils and potential downsides.

At its heart, Gen-AI risk management needs to be proactive and integrated from the ground up. Our own experience with PairD, the Deloitte AI tool, highlights the value of starting with lower-risk AI applications while gradually scaling and refining risk controls[10]. In financial institutions, where sensitive data is handled frequently, stringent security and standardised frameworks are essential for all Gen-AI deployments. Likewise, maintaining full compliance with all relevant data regulations, such as GDPR, and building robust data quality, control and storage strategies are also key lessons.

Trust in AI tools can also be built through transparent decision-making processes and continuous user feedback. Explainability initiatives can help to illuminate the decision-making processes of Gen-AI, and institutions should embrace an iterative approach to risk management, learning from each application and adjusting governance frameworks accordingly to achieve peak compliance and scale with confidence.

For businesses, we would recommend embedding risk management into AI design from the outset, starting with lower risk use cases; prioritising data privacy and ensuring AI systems fully comply with all relevant regulations, including GDPR; and building trust through transparency, continuous monitoring, and user engagement.

How can Government and financial regulators strike the right balance between seizing the opportunities of AI but at the same time protecting consumers and mitigating against any threats to financial stability?

Deloitte has significant experience in conceptualising, developing, and implementing new technologies that support government, commercial, and societal advancement. Whether working with commercial or government clients, this experience has resulted in a proliferation of specialist advice and thought leadership on AI — its promise, but also its limitations.

As highlighted in the AI Opportunities Action Plan this year, *"AI capabilities are developing at an extraordinary pace. If this continues, artificial intelligence (AI) could be the government's single biggest lever to deliver its five missions, especially the goal of kickstarting broad-based economic growth."* We agree it is critical the opportunities of this time are seized, and as set out in the question, an appropriate balance struck.

The Committee will reflect on the full range of options for policymakers to intervene. However, as part of that process, it may wish to draw upon some of the principles, frameworks and guiding principles used by other organisations.

[10] Deloitte. Embracing the Gen-AI revolution – summary recommendations for financial institutions. 2025

We have a range of guiding principles here at Deloitte and can certainly provide further examples and more detail if useful for the Committee. The Deloitte 'Trustworthy AI Framework' developed in 2018 is one of these, and considers the ideal attributes for regulating technologies such as AI. This broad framework proposed core principles with the aim to enhance and enable the most positive attributes of new technologies, while protecting individuals, businesses, and societies from its potentially negative implications.

In addition, Deloitte as an organisation has published extensively on AI[11]. Three recurring themes are highlighted in recent papers. The first theme is the role of AI in advancing fairness across society. It is crucial for those involved in designing, monitoring, utilising, and regulating AI to proactively address bias to avoid perpetuating existing inequities or creating new ones. Policymakers are actively engaged in debates on how to regulate algorithms, assess bias, and promote fairness, with European countries being at the forefront of proposing regulations in this area.

The second theme revolves around trust in AI and its integration into our daily lives. Trustworthy AI, which is ethical, lawful, and technically robust, is necessary to gain broad trust in technology, as it is integrated into everyday technology, whether users are aware of it or not. Policy implications include enhancing algorithm transparency and accountability to help individuals understand AI's processes and conclusions. While organisational policies play a significant role in building trust, national governments also need to consider trust-building measures when deploying AI for providing services to their citizens.

The third theme highlights the role of AI in driving economic prosperity. When implemented effectively, AI can automate tasks and create opportunities for higher-value skilled work. To realise this potential, continued investment in innovation and AI-related technologies is crucial, both at the national level and within organisations.

Last year, the previous UK Government unveiled its long-awaited response to last year's White Paper consultation on regulating Artificial Intelligence (AI). With a new government, evolving technology and ever-changing geopolitical environment, there will be changes in emphasis and updates to policy positions.[12]

It remains clear that organisations must prepare for increased regulatory activity over AI in future including guidelines, information gathering, and enforcement. International firms will also inevitably have to navigate regulatory divergence.

As a firm focused on supporting organisations navigate these changing landscapes, we hold significant expertise in these areas. Our submission to this inquiry has highlighted numerous principles which we view as important, alongside specific use cases. There is a positive story to tell on the use of GenAI in financial services, but as we have noted, there are also risks and uncertainties to manage. To strike an effective balance between these, detail will be key.

We would be happy to support the Committee in advising on any technical matters they may wish to explore, drawing upon expertise from our Deloitte AI Institute, the EMEA Centre for Regulatory Strategy, and other parts of the firm.

April 2025

[11] Deloitte. Artificial Intelligence and Rulemaking. 2025
[12] Deloitte. The UKs framework for AI Regulation. 2025